

DDCA Cheatsheet

Binary Numbers

$2^0 = 1$ $2^3 = 8$ $2^6 = 64$ $2^9 = 512$
 $2^1 = 2$ $2^4 = 16$ $2^7 = 128$ $2^{10} = 1024$ Kilo
 $2^2 = 4$ $2^5 = 32$ $2^8 = 256$ $2^{20} = 1,048,576$ Mega

Sign-Magnitude: First Bit as sign-bit

2's complement: $5 \Rightarrow 0101 \Rightarrow 1010 \Rightarrow 1011 = -5$

Boolean Algebra

$T6$ $B \cdot C = C \cdot B$ $T6'$ $B+C = C+B$ Commutativity
 $T7$ $(B \cdot C) \cdot D = B \cdot (C \cdot D)$ $T7'$ $(B+C)+D = B+(C+D)$ Associativity
 $T8$ $(B \cdot C) + (B \cdot D) = B \cdot (C+D)$ $T8'$ $(B+C) \cdot (B+D) = B+(C \cdot D)$ Distributivity
 $T9$ $B \cdot (B+C) = B$ $T9'$ $B+(B \cdot C) = B$ Covering
 $T10$ $(B \cdot C) + (B \cdot \bar{C}) = B$ $T10'$ $(B+C) \cdot (B+C) = B+C$ Consensus
 $T11$ $(B \cdot C) + (\bar{B} \cdot D) + (C \cdot D) = B \cdot C + \bar{B} \cdot D$ $T11'$ $(B+C) \cdot (\bar{B}+D) \cdot (C+D) = (B+C) \cdot (\bar{B}+D)$ Consensus
 $T12$ $\overline{B_1 \cdot B_2 \cdot \dots} = \bar{B}_1 + \bar{B}_2 + \bar{B}_3 + \dots$ $T12'$ $\overline{B_1 + B_2 + \dots} = \bar{B}_1 \cdot \bar{B}_2 \cdot \bar{B}_3 \cdot \dots$ De Morgan's Theorem

Product of Sum:

A	B	X
0	0	1
1	0	0
0	1	1
1	1	0

$X = (\bar{A} + \bar{B}) \cdot (\bar{A} + B)$
 $X = (\bar{A} \cdot \bar{B}) + (\bar{A} \cdot B)$

Completeness

	Nand	Nor
Not	$\overline{A \cdot A}$	$\overline{A + A}$
And	$\overline{\overline{A \cdot B}}$	$\overline{\overline{A + B}}$
Nand	$\overline{A \cdot B}$	$\overline{A + B}$
Or	$\overline{\overline{A \cdot B}}$	$\overline{\overline{A + B}}$
Nor	$\overline{A \cdot B}$	$\overline{A + B}$
XOR	$\overline{A \cdot B} + \overline{\bar{A} \cdot \bar{B}}$	$\overline{A + B} + \overline{\bar{A} + \bar{B}}$
XNOR	$\overline{A \cdot B} + \bar{A} \cdot \bar{B}$	$\overline{A + B} + A \cdot B$

Karnaugh Maps:

1. Rowest possible 2. size $2^n \times 2^n$ 3. only 1
 4. X may be used as 1's 5. as big as possible


A	B	C	X
0	0	0	1
0	0	1	1
0	1	0	0
0	1	1	0
1	0	0	1
1	0	1	1
1	1	0	1
1	1	1	0

$\Rightarrow X = \bar{A} \cdot \bar{B} + \bar{A} \cdot B + A \cdot \bar{B} + A \cdot B$
 $\Rightarrow X = \bar{A} + B + A$

Big vs. Little Endian:

	0	1	2	3
Big Endian:	ca	fe	2b	3a
Little Endian:	3a	2b	fe	ca

Logic Gates



Verilog Modules

D Flip-Flop
 always @(posedge clk) begin
 if (!rst) q <= 0;
 else q <= d;
 end

MUX
 assign y = s[s[1]?(s[0]?x:y):
 (s[0]?w:a);
 // could be in always block with
 always @(*) begin
 case (var)
 2'b01: out = val;
 2'b10: out = val;
 default: out = val;
 end case
 end

D-Latch
 always @(clk, d) begin
 if (clk) q <= d;
 end

Full Adder
 assign sum = a^b^c-in;
 assign cout = (a&b)|(a&c)|(b&c-in);

Correct Code *twice*: left of assign is no assignment in always-block/connect in/out of modules
 Regs: can't be connected to output port of module/cannot be used in input port declaration/left of <=/
 1/0: check if names match and if all ports are assigned
 General: not multiple assignments to same signal/no recursion/names const stat with numbers

Verilog Operators

!: logic negation ~!: red NAND
 ~: bitwise negation ~!: red NOR
 &: red and ^: red XOR
 |: red or ~^: red XNOR

Combinational vs. Sequential

Comb: - all left-hand signals are assigned in every possible way
 - all inputs are in sensitivity list
 - all outputs are assigned
 - no memory
 - no cyclic paths
 - combined inputs to get output

Seq: - has memory
 - depends on prior inputs
 - not all outputs are assigned in all cases

Finite State Machines

Moore: - Output depends on current state. - more number of states
 - synchronous output & state generation - output placed on states
 Mealy: - Output depends on current state and input
 - less states - output placed on transitions

Designing a FSM

1. Identify inputs & outputs
2. Sketch State Transition diagram
3. Write state transition & output table
4. Write boolean equations for next state

Area of FSM

#FF = #bits for state X 2 - #logic gates = count next state/output logic

State Encodings

Binary Encoding (00, 01, 10, 11)
 (log₂(#states)) bits needed
 reduces # FF to hold states

One-hot Encoding (001, 010, 100)
 #states bits needed
 reduces next state logic

Output Encoding (100, 110, 111)
 reduces output logic

Correctness of FSM

- reset line
 - not multiple transitions for some input
 - no mixing transitions
 - no unmarked transitions
 - initial state
 - no mix of Moore/Mealy

MIPS

R-Type: Register Type, two source/one destination register
I-Type: Immediate Type, one source/one destination register + imm. value
J-Type: Jump Type, operand + address (+ branch)
Reserved
 saved register \$s0 - \$s7
 return address \$ra
 stack pointer \$sp
 stack above pointer
Non-Reserved
 temporary register \$t0 - \$t9
 argument register \$a0 - \$a3
 return value register \$v0 - \$v1
 stack below the stack pointer

Memory Map

↓ Stack	
Dynamic Data	
Heap ↑	
Global Data	
Text	
Reserved	
Reserved	

$\$JIP = 0x7FFFFFFC$ →
accessed with 76-bit positive or neg. immediate
↓
 $\$GP = 0x70000000$ →
 $PC = 0x00000000$ →

can interleave \Rightarrow corruption
Global variables, defined before startup of program
256 MB of code
4 MB are 0, then the j-init can jump to any line of code

ISA

Interface between SW and HW
"what programmer sees"

- Instruction: opcode, addressing mode, data types, instruction type and format, registers, condition codes
- Memory: address space, alignment, addressability, virtual memory management
- Call, interrupt and exception handling
- I/O: memory mapped vs. instruction
- Power & Thermal management
- Multiprocessing / Multithreading support
- Access Control, priority and privilege
- Memory location of exception vectors
- Function of each bit in a programmable branch predictor register
- Order of loads and stores in multi-core CPU
- Program counter width
- Hardware FP-exception support
- Vector instruction support
- CPU endianness
- Virtual Page size

μ-Arch

Specifies underlying implementation that actually executes instructions

- Pipelining
- In order vs. Out-of-Order execution
- Memory address scheduling policy
- Speculative execution
- Superscalar processing
- Clock gating
- Caching: level, size, associativity, replacement policies
- Error correction
- Physical structure
- Instruction latency
- Physical memory page size
- Instruction issue width
- reservation stage capacity
- # pipeline stages
- latency of branch miss prediction
- Fetch width of superscalar CPU
- # non-programmable CPU registers
- register file has one input and two output ports
- number of read ports in physical register file

Performance Evaluation

CPI: cycles per instruction

MIPS: million instructions/sec = MHz/CPI

IPC: instructions per cycle

Time: #instr. \cdot CPI $\cdot \frac{1}{\text{MHz}}$

MHz: frequency, 10^6 cycles/s

Speedup: oldTime / newTime

higher MHz \nRightarrow higher MIPS, IPC could be lower

higher MIPS \nRightarrow less time, could need more instructions

Single Cycle Machines

Each instruction takes a single clock cycle and all state updates are made at the end of the cycle. - slowest instruction determines cycle time

+ easy to build


Multi Cycle Machines

Instructions processing is broken into multiple stages/cycles. State changes happen during execution and architectural updates at the end. Instruction processing consists of two components: • Datapath - relay and transform data • Control logic - FSM for signals

+ slowest stage determines cycle time

Datapath A program consists of **datapath nodes**. A node fires (executes) when all inputs are ready

Datapath Modules



Pipelining

The idea is to process multiple instructions at once by keeping each stage occupied. In reality there are a few problems, which cause pipeline stalls:

- **Data/Control Flow dependencies**: Flow (read after write), output (write after write) and anti (write after read) dependencies. The last two exist due to lack of register file.
- **Resource contention**: can be fixed by duplication, increased throughput or deletion and stalling
- **Long latency operations**

e.g. fine-grained multithreading

Handling flow dependencies

- stall
- eliminate at software level
- predict values
- do something else
- **data forwarding** \rightarrow W \rightarrow D: internal/register file forwarding
- \hookrightarrow M \rightarrow E: operand forwarding

Pipeline stages

- Fetch**: Read instruction from memory
- Decode**: Read operands from register file
- Execute**: ALU-operation
- Memory**: read/write from/to memory
- Writeback**: write result to register file

Interlocking & Scoreboarding

- Detection of data dependencies to ensure correct execution
- SW-Interlocking**: Compiler inserts nops \Rightarrow nops go through all pipeline stages
- HW-Interlocking**: Stall the pipeline

Out of Order Execution

Move dependent instructions out of the way of independent ones.

In-Order Pipeline with Reorder Buffer \rightarrow if yes, dispatch to ALU

- Decode**: Access regfile/ROB, allocate entry in ROB, check if it can execute
- Execute**: Instructions out of order
- Completion**: Write result to reorder buffer \rightarrow else flush pipeline
- Retirement/Commit**: Check exceptions; if none \Rightarrow write architectural register file or Mem.
- \hookrightarrow In-order dispatch/execution, OoO completion, in-order retirement.

Tomasulo's Algorithm

Implementing OoO execution. Uses register renaming to eliminate output and anti-dependencies. It rather uses reservation stations for individual ops.

1. If reservation station is available \rightarrow comes from register alias table RAT
 - instr. + renamed operands inserted into reservation station
 - rename destination register in RAT
 - Else: Stall pipeline
2. While in reservation station:
 - watch common data bus for tag of source
 - if tag seen \Rightarrow grab value \Rightarrow set valid bit
 - if both operands are valid \Rightarrow instr. ready for dispatch
3. Dispatch instruction to functional unit
4. After instruction finishes
 - put tagged value onto common data bus
 - if register alias table contain tag \Rightarrow update value and set valid bit \hookrightarrow write to register file
 - redoin rename tag \rightarrow no valid copy of tag in the system

VLIW

Compiler finds independent instructions and schedule them into a single VLIW-instr.

- lock step execution**: if one instruction stalls, the whole VLIW stalls
- + simple hardware
- compiler needs to find N independent instructions
- + no dependency checking
- lock step cause stalls \hookrightarrow complex
- + no instruction distribution

Superscalar Execution

- Fetch/Decode/... multiple instructions per cycle
- + higher IPC
- higher complexity for dependency checking \Rightarrow more HW

Systolic Arrays

Instead of a single processing element (PE), we have an array of PE and carefully orchestrate the datapath between them. \Rightarrow max. comp. on single PE


Difference to Pipelining: Array structure is non-linear and multi-dimensional. PE-structure can be multi-directional on different speeds. PE can have memory.

Fine Grained Multithreading

HW has multiple thread contexts (PC + reg) and each cycle the fetch-engine fetches from a different thread.


- + no dependencies
- extra hardware
- + no branch prediction
- reduced single-thread performance
- + improved throughput, latency, tolerance, utilization
- resource contention \hookrightarrow dependency checking between threads

SIMD Single Instruction operates on multiple data



Memory & Memory Hierarchy

Memory Arrays: stores data, address selection, logic select row, readout circuitry, reads data



Flip-Flop (latches)

+ Very fast, parallel access

- Very expensive

Static RAM

+ fast

- serial, expensive

Dynamic RAM

+ cheap

- volatile, slow, serial

Locality: temporal = access to same address in short-time
 spatial = access to nearby address


Caching (Exploit locality) Store adjacent data/recently accessed data

Blocks & Addressing cache:

• Memory is divided into fixed-size blocks

• each block/address maps to a set in the cache

• For a cache hit, the tag needs to match + valid bit



Associativity

• n-way associative cache: n blocks per set/allows n blocks with same index

• fully associative cache: 1 set \equiv 0 index bits/Memory address can map to any block

• direct mapped cache: 1 block per set

Cache Performance

• cache size, total data = c

• block size = b

• associativity = n

Misrate = # Misses / # total Mem. accesses

Hitrate = # Hits / # total Mem. accesses

Blocks: $B = \frac{c}{b}$

sets: $S = \frac{B}{n} = \frac{c}{b \cdot n}$

Replacement Policies

• FIFO (first written)

• LRU (Least accessed)

• random

Write Policy

Writeback: dirty-bit/write to lower levels when block is evicted

Write-through: write to all levels immediately \Rightarrow simpler, but bandwidth intensive

Types of misses

• compulsory miss: first request to cache is always a miss

• capacity miss: cache too small to hold all concurrently used data

• conflict miss: several addresses map to the same set and evict still needed blocks

Improve Cache Performance

Reduce miss rate:

- more associativity
- better replacement policies
- software approaches
- prefetching

Reduce miss latency/cost:

- multi-level caches
- critical words first
- better replacement policy
- software approaches
- prefetching

Prefetching Preload data to avoid misses in cache. Done by SW/HW

stride prefetcher: prefetches cache block in a pattern with certain stride

\hookrightarrow if stride = 0: next-block prefetching / n-block prefetcher

Runahead execution: allows the processor to pre-process instructions during cache miss instead of stalling. Therefore it can detect potential cache misses earlier.

Performance coverage = #correctly prefetched / #accessed blocks

accuracy = #correctly prefetched / #total prefetched blocks

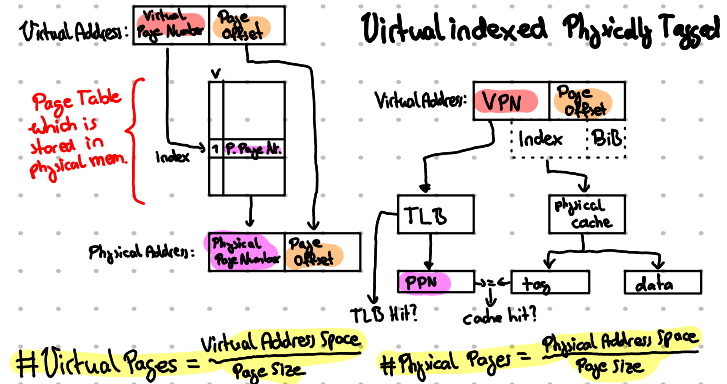
Virtual Memory

Much larger than physical memory. Virtual address space divided into pages

Physical address space divided into frames. Page Table stores mapping: $V \rightarrow P$


If accessed virtual page is not in memory, but on disk \Rightarrow load in memory (demand paging)

! Physical Memory is a cache for pages stored on disk (fully associative) !



Multi-Level page tables:

Keeps PT size small



Memory Protection: different PT for each program

Translation Lookaside Buffer (TLB): cache PT entries to speed up address translation

Vector Processing: Perform op on a whole array. Only possible if elements are independent

The vector data is stored in N M-bit vector data registers. **Vector Control Registers:**

Vector Chaining: Data forwarding from one vector functional unit to another

Vector Stripmining: #data elements > #elements in vector register \Rightarrow split into smaller vectors

Memory Banking: Memory divided into banks. Access is independent

Can start and complete 1 access per cycle

\Rightarrow can sustain N concurrent accesses if all N are in diff. banks

+ a lot of work per instruction

+ regular memory access pattern

+ no need for loops

- works only if parallelism is regular otherwise it is highly inefficient

GPU SIMD engines but programmed using threads (Not SIMD instructions)

SIMT: SIMD instr. (VADD, ...) not exposed to programmers \Rightarrow treat threads separately

A set of threads executing the same instruction are dynamically grouped into a warp

Dynamic Warp Merging: Dynamically merge threads executing the same instr. after branch divergence

#warps = #threads / #threads per warp

Utilization = #instr. executed / best case when all warp are full instr.

\hookrightarrow all branches are taken the same, if > 1 instr. takes branch

Branch Prediction

A technique used to predict the next address after a branch. If the prediction is wrong, we need to flush the pipeline (misprediction penalty)

Prediction Schemes

• Always (not) taken

• BTBN (Backward taken, forward not taken)

• Profile Based (likely direction) hint from compiler

• last time predictor: single bit in BTB

• 2-bit counter based prediction

• Advanced algorithms (perceptrons)

Local correlation: PC indexes local history registers. Entry used to index PHT

Global correlation: global T/NT history of all branches in GHR (used to index PHT)

PC indexes

local history register

Pattern History Table

Global History Register

Pattern History Table